

DV0101EN-2-2-1-Area-Plots-Histograms-and-Bar-Charts-py-v2.0

March 17, 2019

Area Plots, Histograms, and Bar Plots

0.1 Introduction

In this lab, we will continue exploring the Matplotlib library and will learn how to create additional plots, namely area plots, histograms, and bar charts.

0.2 Table of Contents

1. Section ??
2. Section ??
3. Section ??
4. Section ??
5. Section ??
6. Section ??

1 Exploring Datasets with *pandas* and Matplotlib

Toolkits: The course heavily relies on *pandas* and *Numpy* for data wrangling, analysis, and visualization. The primary plotting library that we are exploring in the course is *Matplotlib*.

Dataset: Immigration to Canada from 1980 to 2013 - [International migration flows to and from selected countries - The 2015 revision](#) from United Nation's website.

The dataset contains annual data on the flows of international migrants as recorded by the countries of destination. The data presents both inflows and outflows according to the place of birth, citizenship or place of previous / next residence both for foreigners and nationals. For this lesson, we will focus on the Canadian Immigration data.

2 Downloading and Prepping Data

Import Primary Modules. The first thing we'll do is import two key data analysis modules: *pandas* and *Numpy*.

```
In [1]: import numpy as np # useful for many scientific computing in Python
import pandas as pd # primary data structure library
```

Let's download and import our primary Canadian Immigration dataset using *pandas* `read_excel()` method. Normally, before we can do that, we would need to download a module which *pandas* requires to read in excel files. This module is **xlrd**. For your convenience, we have pre-installed this module, so you would not have to worry about that. Otherwise, you would need to run the following line of code to install the **xlrd** module:

```
!conda install -c anaconda xlrd --yes
```

Download the dataset and read it into a *pandas* dataframe.

```
In [2]: df_can = pd.read_excel('https://ibm.box.com/shared/static/lw190pt9zpy5bd1ptyg2aw15awomz9
                                sheet_name='Canada by Citizenship',
                                skiprows=range(20),
                                skipfooter=2
                                )

    print('Data downloaded and read into a dataframe!')
```

Data downloaded and read into a dataframe!

Let's take a look at the first five items in our dataset.

```
In [3]: df_can.head()
```

```
Out[3]:
```

	Type	Coverage	OdName	AREA	AreaName	REG	\
0	Immigrants	Foreigners	Afghanistan	935	Asia	5501	
1	Immigrants	Foreigners	Albania	908	Europe	925	
2	Immigrants	Foreigners	Algeria	903	Africa	912	
3	Immigrants	Foreigners	American Samoa	909	Oceania	957	
4	Immigrants	Foreigners	Andorra	908	Europe	925	

	RegName	DEV	DevName	1980	...	2004	2005	2006	\
0	Southern Asia	902	Developing regions	16	...	2978	3436	3009	
1	Southern Europe	901	Developed regions	1	...	1450	1223	856	
2	Northern Africa	902	Developing regions	80	...	3616	3626	4807	
3	Polynesia	902	Developing regions	0	...	0	0	1	
4	Southern Europe	901	Developed regions	0	...	0	0	1	

	2007	2008	2009	2010	2011	2012	2013
0	2652	2111	1746	1758	2203	2635	2004
1	702	560	716	561	539	620	603
2	3623	4005	5393	4752	4325	3774	4331
3	0	0	0	0	0	0	0
4	1	0	0	0	0	1	1

[5 rows x 43 columns]

Let's find out how many entries there are in our dataset.

```
In [4]: # print the dimensions of the dataframe
print(df_can.shape)
```

```
(195, 43)
```

Clean up data. We will make some modifications to the original dataset to make it easier to create our visualizations. Refer to Introduction to Matplotlib and Line Plots lab for the rational and detailed description of the changes.

1. Clean up the dataset to remove columns that are not informative to us for visualization (eg. Type, AREA, REG).

```
In [5]: df_can.drop(['AREA', 'REG', 'DEV', 'Type', 'Coverage'], axis=1, inplace=True)
```

```
# let's view the first five elements and see how the dataframe was changed
df_can.head()
```

```
Out[5]:
```

	OdName	AreaName	RegName	DevName	1980	1981	\
0	Afghanistan	Asia	Southern Asia	Developing regions	16	39	
1	Albania	Europe	Southern Europe	Developed regions	1	0	
2	Algeria	Africa	Northern Africa	Developing regions	80	67	
3	American Samoa	Oceania	Polynesia	Developing regions	0	1	
4	Andorra	Europe	Southern Europe	Developed regions	0	0	

	1982	1983	1984	1985	...	2004	2005	2006	2007	2008	2009	2010	\
0	39	47	71	340	...	2978	3436	3009	2652	2111	1746	1758	
1	0	0	0	0	...	1450	1223	856	702	560	716	561	
2	71	69	63	44	...	3616	3626	4807	3623	4005	5393	4752	
3	0	0	0	0	...	0	0	1	0	0	0	0	
4	0	0	0	0	...	0	0	1	1	0	0	0	

	2011	2012	2013
0	2203	2635	2004
1	539	620	603
2	4325	3774	4331
3	0	0	0
4	0	1	1

```
[5 rows x 38 columns]
```

Notice how the columns Type, Coverage, AREA, REG, and DEV got removed from the dataframe.

2. Rename some of the columns so that they make sense.

```
In [6]: df_can.rename(columns={'OdName':'Country', 'AreaName':'Continent', 'RegName':'Region'}, i
```

```
# let's view the first five elements and see how the dataframe was changed
df_can.head()
```

```

Out[6]:
      Country Continent      Region      DevName  1980  1981  \
0  Afghanistan      Asia  Southern Asia  Developing regions    16    39
1      Albania      Europe  Southern Europe    Developed regions     1     0
2      Algeria      Africa  Northern Africa  Developing regions    80    67
3  American Samoa  Oceania      Polynesia  Developing regions     0     1
4      Andorra      Europe  Southern Europe    Developed regions     0     0

      1982  1983  1984  1985  ...  2004  2005  2006  2007  2008  2009  2010  \
0     39    47    71   340  ...  2978  3436  3009  2652  2111  1746  1758
1      0     0     0     0  ...  1450  1223   856   702   560   716   561
2     71    69    63    44  ...  3616  3626  4807  3623  4005  5393  4752
3      0     0     0     0  ...     0     0     1     0     0     0     0
4      0     0     0     0  ...     0     0     1     1     0     0     0

      2011  2012  2013
0  2203  2635  2004
1   539   620   603
2  4325  3774  4331
3      0     0     0
4      0     1     1

[5 rows x 38 columns]

```

Notice how the column names now make much more sense, even to an outsider.

3. For consistency, ensure that all column labels of type string.

```

In [7]: # let's examine the types of the column labels
      all(isinstance(column, str) for column in df_can.columns)

```

```

Out[7]: False

```

Notice how the above line of code returned *False* when we tested if all the column labels are of type **string**. So let's change them all to **string** type.

```

In [8]: df_can.columns = list(map(str, df_can.columns))

      # let's check the column labels types now
      all(isinstance(column, str) for column in df_can.columns)

```

```

Out[8]: True

```

4. Set the country name as index - useful for quickly looking up countries using .loc method.

```

In [9]: df_can.set_index('Country', inplace=True)

      # let's view the first five elements and see how the dataframe was changed
      df_can.head()

```

```

Out[9]:
      Continent      Region      DevName  1980  1981  \
Country
Afghanistan      Asia      Southern Asia  Developing regions    16    39
Albania           Europe      Southern Europe    Developed regions     1     0
Algeria           Africa      Northern Africa  Developing regions    80    67
American Samoa   Oceania           Polynesia  Developing regions     0     1
Andorra           Europe      Southern Europe    Developed regions     0     0

      1982  1983  1984  1985  1986  ...  2004  2005  2006  2007  \
Country
Afghanistan      39    47    71   340   496  ...  2978  3436  3009  2652
Albania           0     0     0     0     1  ...  1450  1223   856   702
Algeria           71    69    63    44    69  ...  3616  3626  4807  3623
American Samoa    0     0     0     0     0  ...     0     0     1     0
Andorra           0     0     0     0     2  ...     0     0     1     1

      2008  2009  2010  2011  2012  2013
Country
Afghanistan    2111  1746  1758  2203  2635  2004
Albania         560   716   561   539   620   603
Algeria        4005  5393  4752  4325  3774  4331
American Samoa     0     0     0     0     0     0
Andorra          0     0     0     0     1     1

```

[5 rows x 37 columns]

Notice how the country names now serve as indices.

5. Add total column.

```

In [10]: df_can['Total'] = df_can.sum(axis=1)

# let's view the first five elements and see how the dataframe was changed
df_can.head()

```

```

Out[10]:
      Continent      Region      DevName  1980  1981  \
Country
Afghanistan      Asia      Southern Asia  Developing regions    16    39
Albania           Europe      Southern Europe    Developed regions     1     0
Algeria           Africa      Northern Africa  Developing regions    80    67
American Samoa   Oceania           Polynesia  Developing regions     0     1
Andorra           Europe      Southern Europe    Developed regions     0     0

      1982  1983  1984  1985  1986  ...  2005  2006  2007  2008  \
Country
Afghanistan      39    47    71   340   496  ...  3436  3009  2652  2111
Albania           0     0     0     0     1  ...  1223   856   702   560
Algeria           71    69    63    44    69  ...  3626  4807  3623  4005

```

American Samoa	0	0	0	0	0	...	0	1	0	0
Andorra	0	0	0	0	2	...	0	1	1	0

	2009	2010	2011	2012	2013	Total
Country						
Afghanistan	1746	1758	2203	2635	2004	58639
Albania	716	561	539	620	603	15699
Algeria	5393	4752	4325	3774	4331	69439
American Samoa	0	0	0	0	0	6
Andorra	0	0	0	1	1	15

[5 rows x 38 columns]

Now the dataframe has an extra column that presents the total number of immigrants from each country in the dataset from 1980 - 2013. So if we print the dimension of the data, we get:

```
In [11]: print ('data dimensions:', df_can.shape)
```

```
data dimensions: (195, 38)
```

So now our dataframe has 38 columns instead of 37 columns that we had before.

```
In [12]: # finally, let's create a list of years from 1980 - 2013
         # this will come in handy when we start plotting the data
         years = list(map(str, range(1980, 2014)))
```

years

```
Out[12]: ['1980',
          '1981',
          '1982',
          '1983',
          '1984',
          '1985',
          '1986',
          '1987',
          '1988',
          '1989',
          '1990',
          '1991',
          '1992',
          '1993',
          '1994',
          '1995',
          '1996',
          '1997',
          '1998',
          '1999',
```

```
'2000',  
'2001',  
'2002',  
'2003',  
'2004',  
'2005',  
'2006',  
'2007',  
'2008',  
'2009',  
'2010',  
'2011',  
'2012',  
'2013']
```

3 Visualizing Data using Matplotlib

Import Matplotlib and Numpy.

```
In [13]: # use the inline backend to generate the plots within the browser  
         %matplotlib inline  
  
         import matplotlib as mpl  
         import matplotlib.pyplot as plt  
  
         mpl.style.use('ggplot') # optional: for ggplot-like style  
  
         # check for latest version of Matplotlib  
         print ('Matplotlib version: ', mpl.__version__) # >= 2.0.0
```

Matplotlib version: 3.0.2

4 Area Plots

In the last module, we created a line plot that visualized the top 5 countries that contributed the most immigrants to Canada from 1980 to 2013. With a little modification to the code, we can visualize this plot as a cumulative plot, also known as a **Stacked Line Plot** or **Area plot**.

```
In [14]: df_can.sort_values(['Total'], ascending=False, axis=0, inplace=True)  
  
         # get the top 5 entries  
         df_top5 = df_can.head()  
  
         # transpose the dataframe  
         df_top5 = df_top5[years].transpose()  
  
         df_top5.head()
```

```
Out[14]: Country India China United Kingdom of Great Britain and Northern Ireland \
1980      8880   5123                                     22045
1981      8670   6682                                     24796
1982      8147   3308                                     20620
1983      7338   1863                                     10015
1984      5704   1527                                     10170

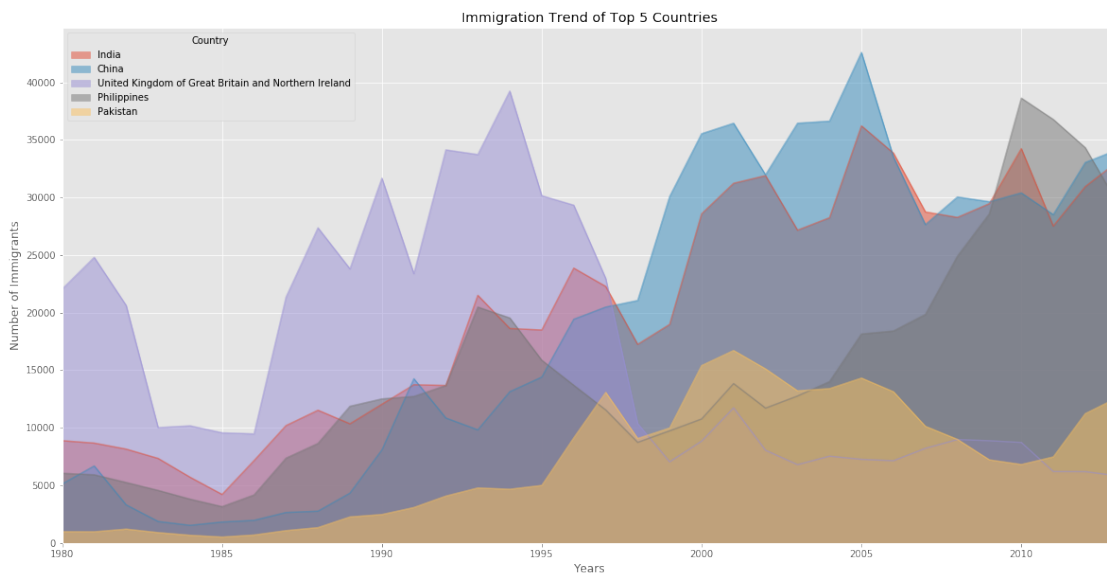
Country Philippines Pakistan
1980              6051      978
1981              5921      972
1982              5249     1201
1983              4562      900
1984              3801      668
```

Area plots are stacked by default. And to produce a stacked area plot, each column must be either all positive or all negative values (any NaN values will default to 0). To produce an unstacked plot, pass `stacked=False`.

```
In [15]: df_top5.index = df_top5.index.map(int) # let's change the index values of df_top5 to type int
df_top5.plot(kind='area',
              stacked=False,
              figsize=(20, 10), # pass a tuple (x, y) size
              )
```

```
plt.title('Immigration Trend of Top 5 Countries')
plt.ylabel('Number of Immigrants')
plt.xlabel('Years')
```

```
plt.show()
```

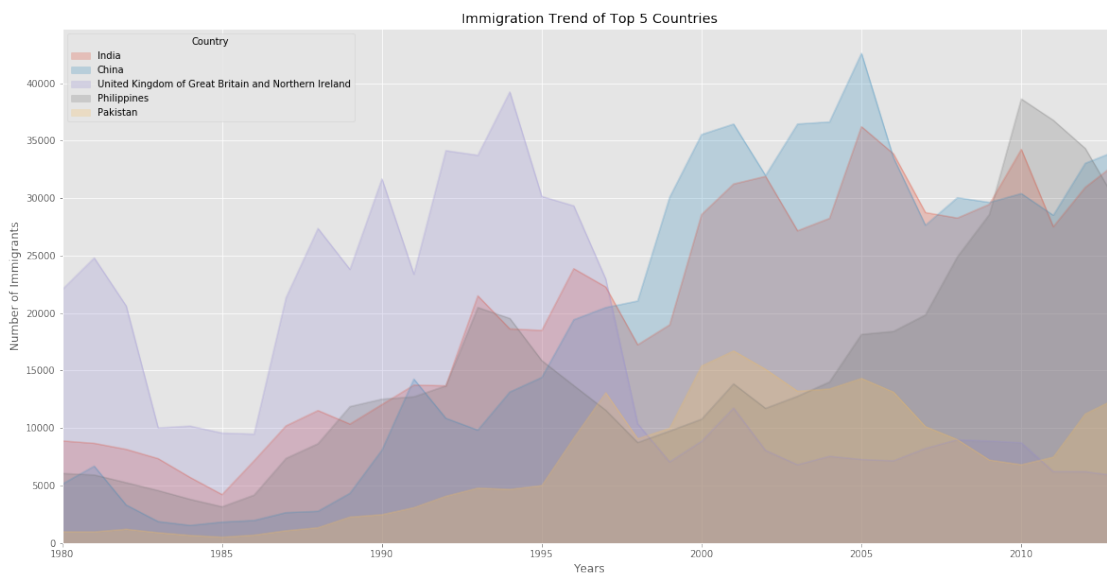


The unstacked plot has a default transparency (alpha value) at 0.5. We can modify this value by passing in the alpha parameter.

```
In [16]: df_top5.plot(kind='area',  
                    alpha=0.25, # 0-1, default value a= 0.5  
                    stacked=False,  
                    figsize=(20, 10),  
                    )
```

```
plt.title('Immigration Trend of Top 5 Countries')  
plt.ylabel('Number of Immigrants')  
plt.xlabel('Years')
```

```
plt.show()
```



4.0.1 Two types of plotting

As we discussed in the video lectures, there are two styles/options of plotting with matplotlib. Plotting using the Artist layer and plotting using the scripting layer.

Option 1: Scripting layer (procedural method) - using matplotlib.pyplot as 'plt'

You can use plt i.e. matplotlib.pyplot and add more elements by calling different methods procedurally; for example, plt.title(...) to add title or plt.xlabel(...) to add label to the x-axis.

```
# Option 1: This is what we have been using so far  
df_top5.plot(kind='area', alpha=0.35, figsize=(20, 10))  
plt.title('Immigration trend of top 5 countries')  
plt.ylabel('Number of immigrants')  
plt.xlabel('Years')
```

Option 2: Artist layer (Object oriented method) - using an `Axes` instance from Matplotlib (preferred)

You can use an `Axes` instance of your current plot and store it in a variable (eg. `ax`). You can add more elements by calling methods with a little change in syntax (by adding "`set_`" to the previous methods). For example, use `ax.set_title()` instead of `plt.title()` to add title, or `ax.set_xlabel()` instead of `plt.xlabel()` to add label to the x-axis.

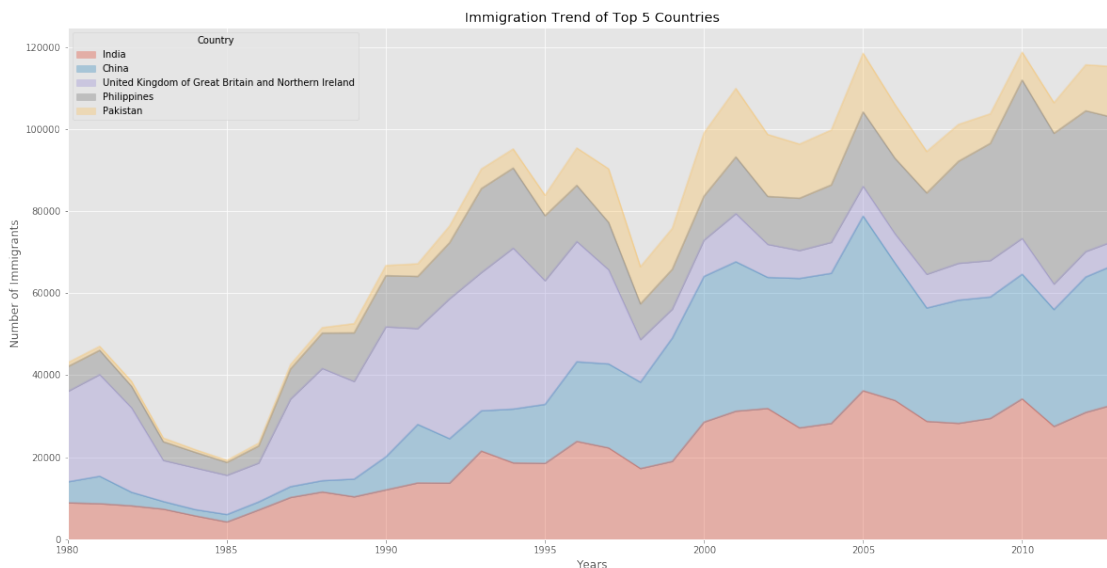
This option sometimes is more transparent and flexible to use for advanced plots (in particular when having multiple plots, as you will see later).

In this course, we will stick to the **scripting layer**, except for some advanced visualizations where we will need to use the **artist layer** to manipulate advanced aspects of the plots.

```
In [17]: # option 2: preferred option with more flexibility
ax = df_top5.plot(kind='area', alpha=0.35, figsize=(20, 10))

ax.set_title('Immigration Trend of Top 5 Countries')
ax.set_ylabel('Number of Immigrants')
ax.set_xlabel('Years')
```

```
Out[17]: Text(0.5, 0, 'Years')
```



Question: Use the scripting layer to create a stacked area plot of the 5 countries that contributed the least to immigration to Canada **from** 1980 to 2013. Use a transparency value of 0.45.

```
In [19]: ### type your answer here
df_least5 = df_can.tail(5)

# transpose the dataframe
df_least5 = df_least5[years].transpose()
df_least5.head()
```

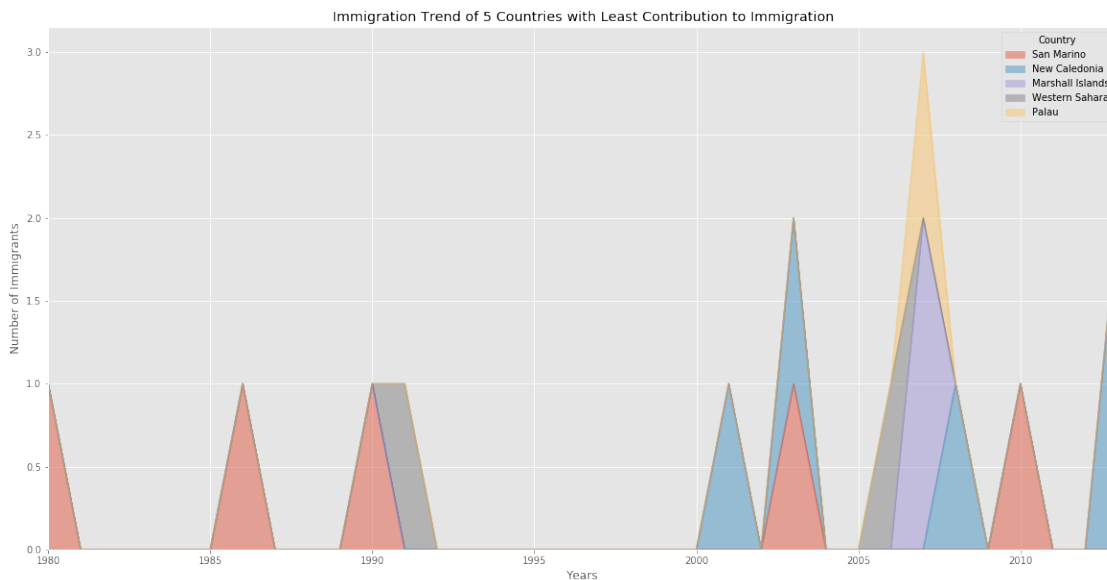
```

df_least5.index = df_least5.index.map(int) # let's change the index values of df_least5
df_least5.plot(kind='area', alpha=0.45, figsize=(20, 10))

plt.title('Immigration Trend of 5 Countries with Least Contribution to Immigration')
plt.ylabel('Number of Immigrants')
plt.xlabel('Years')

plt.show()

```



Question: Use the artist layer to create an unstacked area plot of the 5 countries that contributed the least to immigration to Canada **from** 1980 to 2013. Use a transparency value of 0.55.

In [20]: *### type your answer here*

```

df_least5 = df_can.tail(5)

# transpose the dataframe
df_least5 = df_least5[years].transpose()
df_least5.head()

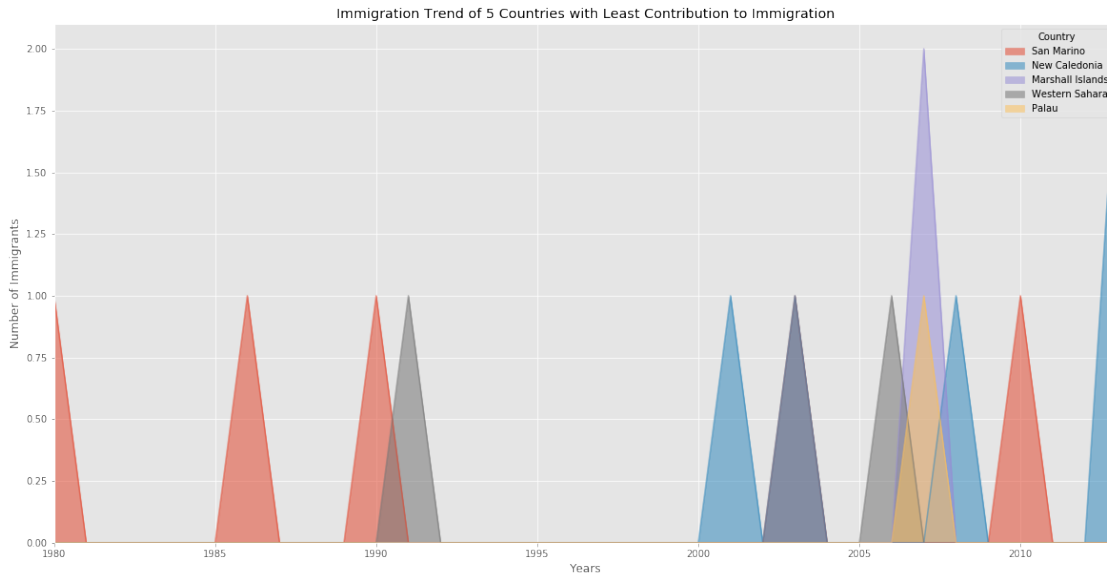
df_least5.index = df_least5.index.map(int) # let's change the index values of df_least5

ax = df_least5.plot(kind='area', alpha=0.55, stacked=False, figsize=(20, 10))

ax.set_title('Immigration Trend of 5 Countries with Least Contribution to Immigration')
ax.set_ylabel('Number of Immigrants')
ax.set_xlabel('Years')

```

Out[20]: Text(0.5, 0, 'Years')



5 Histograms

A histogram is a way of representing the *frequency* distribution of numeric dataset. The way it works is it partitions the x-axis into *bins*, assigns each data point in our dataset to a bin, and then counts the number of data points that have been assigned to each bin. So the y-axis is the frequency or the number of data points in each bin. Note that we can change the bin size and usually one needs to tweak it so that the distribution is displayed nicely.

Question: What is the frequency distribution of the number (population) of new immigrants from the various countries to Canada in 2013?

Before we proceed with creating the histogram plot, let's first examine the data split into intervals. To do this, we will use **Numpy's** histogram method to get the bin ranges and frequency counts as follows:

```
In [21]: # let's quickly view the 2013 data
         df_can['2013'].head()
```

```
Out[21]: Country
         India                33087
         China                34129
         United Kingdom of Great Britain and Northern Ireland    5827
         Philippines          29544
         Pakistan            12603
         Name: 2013, dtype: int64
```

```
In [23]: # np.histogram returns 2 values
         count, bin_edges = np.histogram(df_can['2013'])

         print(count) # frequency count
         print(bin_edges) # bin ranges, default = 10 bins
```

```
[178  11   1   2   0   0   0   0   1   2]
[    0.  3412.9 6825.8 10238.7 13651.6 17064.5 20477.4 23890.3 27303.2
 30716.1 34129. ]
```

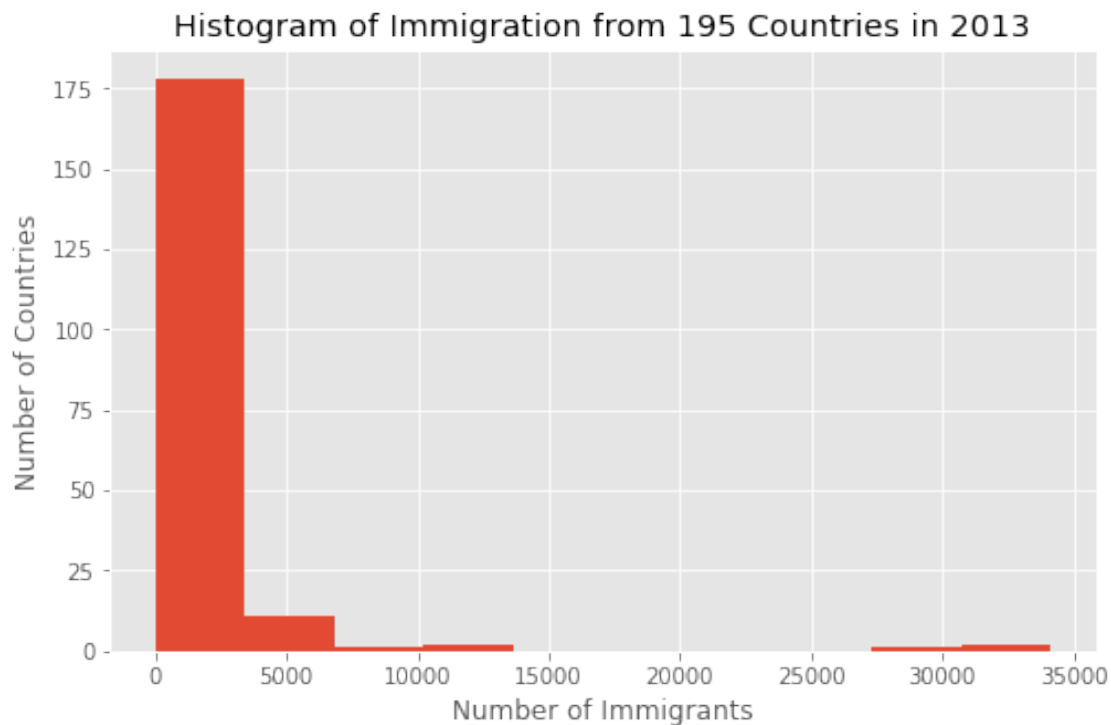
By default, the `histogram` method breaks up the dataset into 10 bins. The figure below summarizes the bin ranges and the frequency distribution of immigration in 2013. We can see that in 2013: * 178 countries contributed between 0 to 3412.9 immigrants * 11 countries contributed between 3412.9 to 6825.8 immigrants * 1 country contributed between 6825.8 to 10238.7 immigrants, and so on..

We can easily graph this distribution by passing `kind=hist` to `plot()`.

```
In [24]: df_can['2013'].plot(kind='hist', figsize=(8, 5))
```

```
plt.title('Histogram of Immigration from 195 Countries in 2013') # add a title to the h
plt.ylabel('Number of Countries') # add y-label
plt.xlabel('Number of Immigrants') # add x-label

plt.show()
```



In the above plot, the x-axis represents the population range of immigrants in intervals of 3412.9. The y-axis represents the number of countries that contributed to the aforementioned population.

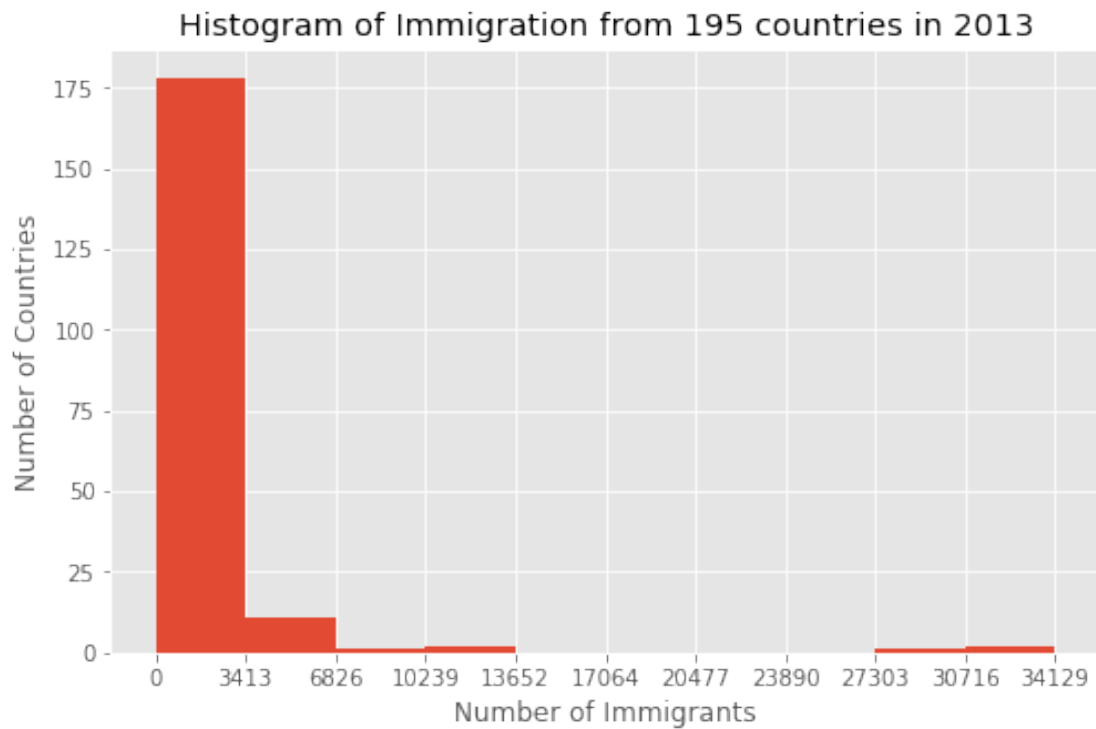
Notice that the x-axis labels do not match with the bin size. This can be fixed by passing in a `xticks` keyword that contains the list of the bin sizes, as follows:

```
In [25]: # 'bin_edges' is a list of bin intervals
count, bin_edges = np.histogram(df_can['2013'])

df_can['2013'].plot(kind='hist', figsize=(8, 5), xticks=bin_edges)

plt.title('Histogram of Immigration from 195 countries in 2013') # add a title to the h
plt.ylabel('Number of Countries') # add y-label
plt.xlabel('Number of Immigrants') # add x-label

plt.show()
```



Side Note: We could use `df_can['2013'].plot.hist()`, instead. In fact, throughout this lesson, using `some_data.plot(kind='type_plot', ...)` is equivalent to `some_data.plot.type_plot(...)`. That is, passing the type of the plot as argument or method behaves the same.

See the *pandas* documentation for more info <http://pandas.pydata.org/pandas-docs/stable/generated/pandas.Series.plot.html>.

We can also plot multiple histograms on the same plot. For example, let's try to answer the following questions using a histogram.

Question: What is the immigration distribution for Denmark, Norway, and Sweden for years 1980 - 2013?

```
In [26]: # let's quickly view the dataset
df_can.loc[['Denmark', 'Norway', 'Sweden'], years]
```

```
Out[26]:
```

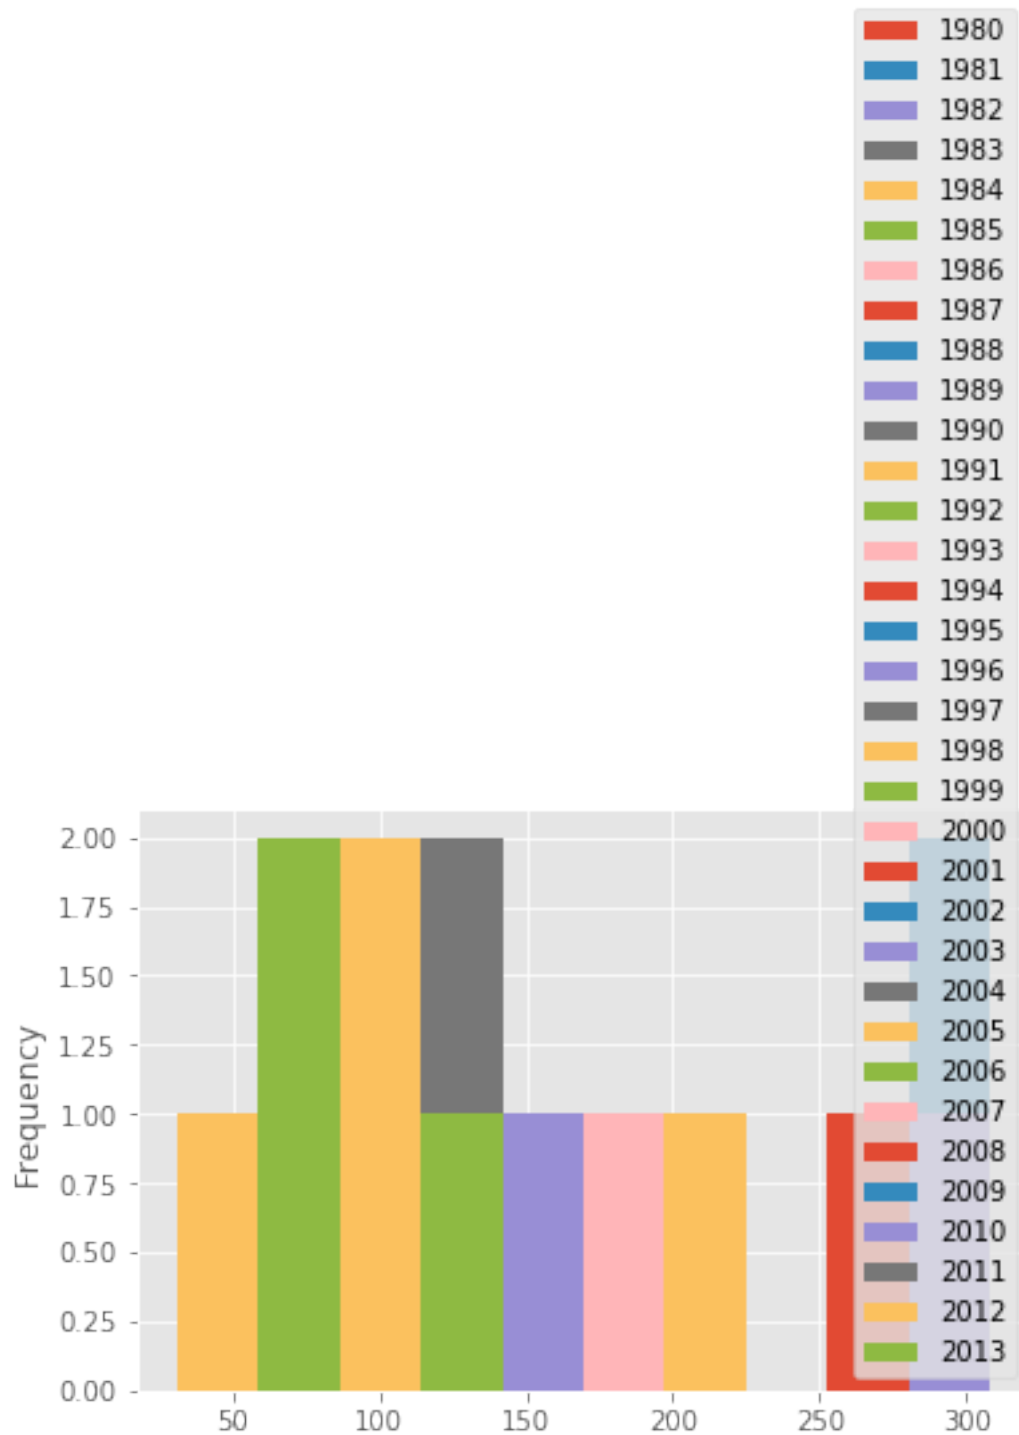
	1980	1981	1982	1983	1984	1985	1986	1987	1988	1989	...	\
Country											...	
Denmark	272	293	299	106	93	73	93	109	129	129	...	
Norway	116	77	106	51	31	54	56	80	73	76	...	
Sweden	281	308	222	176	128	158	187	198	171	182	...	

	2004	2005	2006	2007	2008	2009	2010	2011	2012	2013
Country										
Denmark	89	62	101	97	108	81	92	93	94	81
Norway	73	57	53	73	66	75	46	49	53	59
Sweden	129	205	139	193	165	167	159	134	140	140

[3 rows x 34 columns]

```
In [27]: # generate histogram
df_can.loc[['Denmark', 'Norway', 'Sweden'], years].plot.hist()
```

```
Out[27]: <matplotlib.axes._subplots.AxesSubplot at 0x7f1ed4ad9940>
```



That does not look right!

Don't worry, you'll often come across situations like this when creating plots. The solution often lies in how the underlying dataset is structured.

Instead of plotting the population frequency distribution of the population for the 3 countries,

pandas instead plotted the population frequency distribution for the years.

This can be easily fixed by first transposing the dataset, and then plotting as shown below.

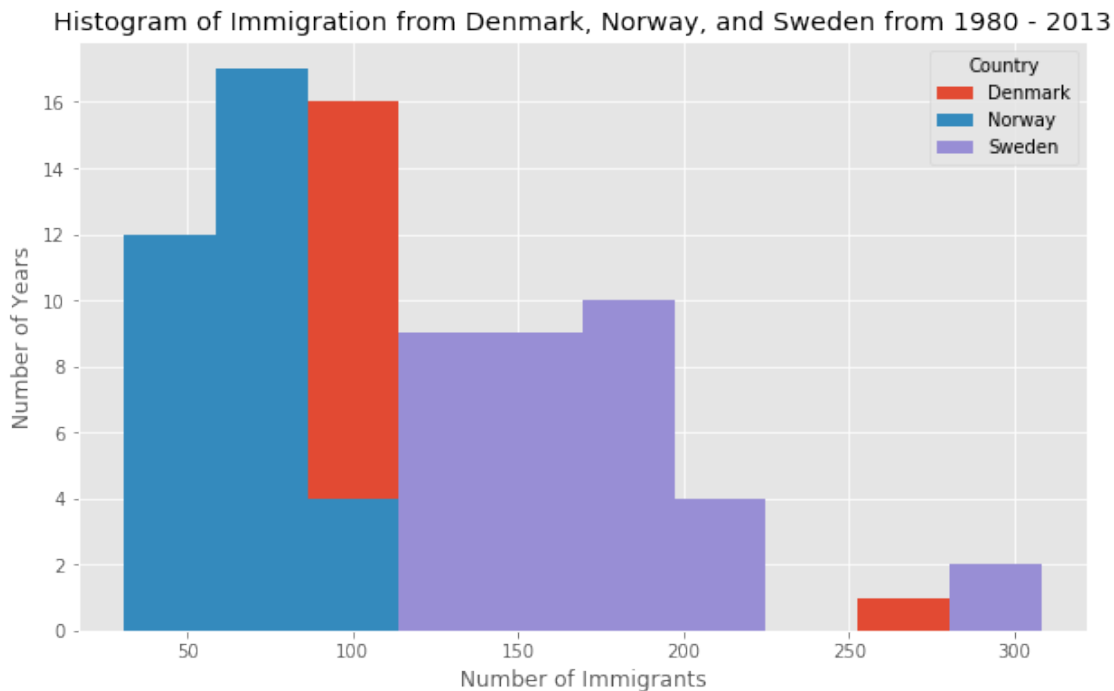
```
In [28]: # transpose dataframe
df_t = df_can.loc[['Denmark', 'Norway', 'Sweden'], years].transpose()
df_t.head()
```

```
Out[28]: Country  Denmark  Norway  Sweden
1980           272      116      281
1981           293       77      308
1982           299      106      222
1983           106       51      176
1984           93       31      128
```

```
In [29]: # generate histogram
df_t.plot(kind='hist', figsize=(10, 6))

plt.title('Histogram of Immigration from Denmark, Norway, and Sweden from 1980 - 2013')
plt.ylabel('Number of Years')
plt.xlabel('Number of Immigrants')

plt.show()
```



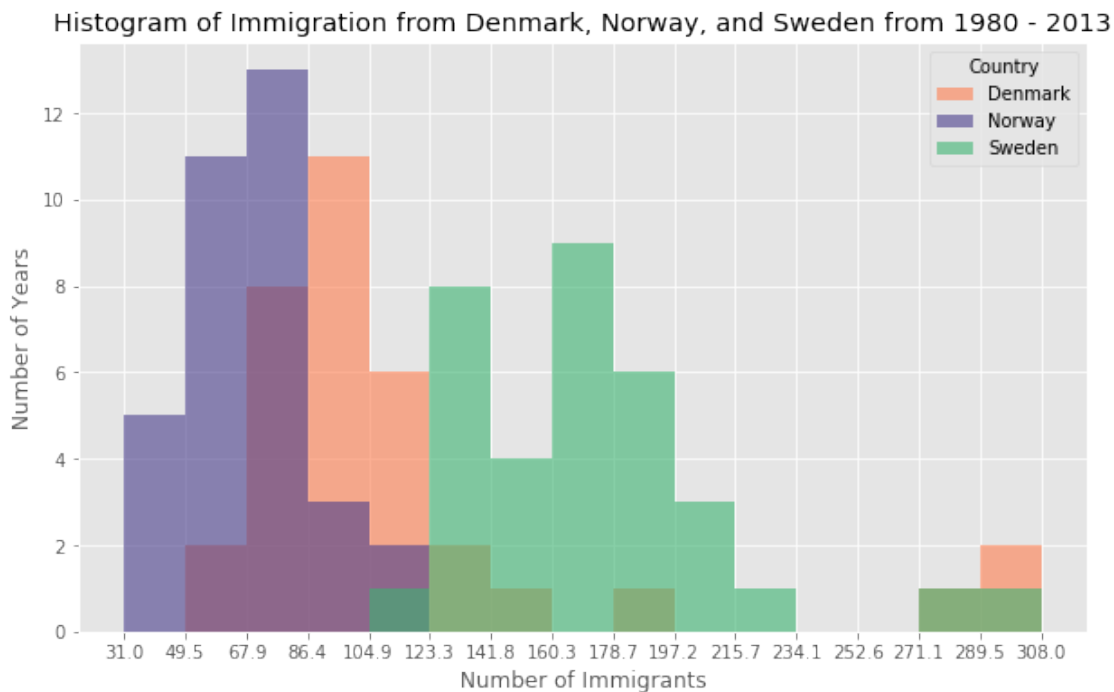
Let's make a few modifications to improve the impact and aesthetics of the previous plot: * increase the bin size to 15 by passing in bins parameter * set transparency to 60% by passing in alpha parameter * label the x-axis by passing in x-label parameter * change the colors of the plots by passing in color parameter

```
In [30]: # let's get the x-tick values
count, bin_edges = np.histogram(df_t, 15)

# un-stacked histogram
df_t.plot(kind='hist',
          figsize=(10, 6),
          bins=15,
          alpha=0.6,
          xticks=bin_edges,
          color=['coral', 'darkslateblue', 'mediumseagreen']
        )

plt.title('Histogram of Immigration from Denmark, Norway, and Sweden from 1980 - 2013')
plt.ylabel('Number of Years')
plt.xlabel('Number of Immigrants')

plt.show()
```



Tip: For a full listing of colors available in Matplotlib, run the following code in your python shell:

```
import matplotlib
for name, hex in matplotlib.colors.cnames.items():
    print(name, hex)
```

If we do not want the plots to overlap each other, we can stack them using the stacked parameter. Let's also adjust the min and max x-axis labels to remove the extra gap on the edges

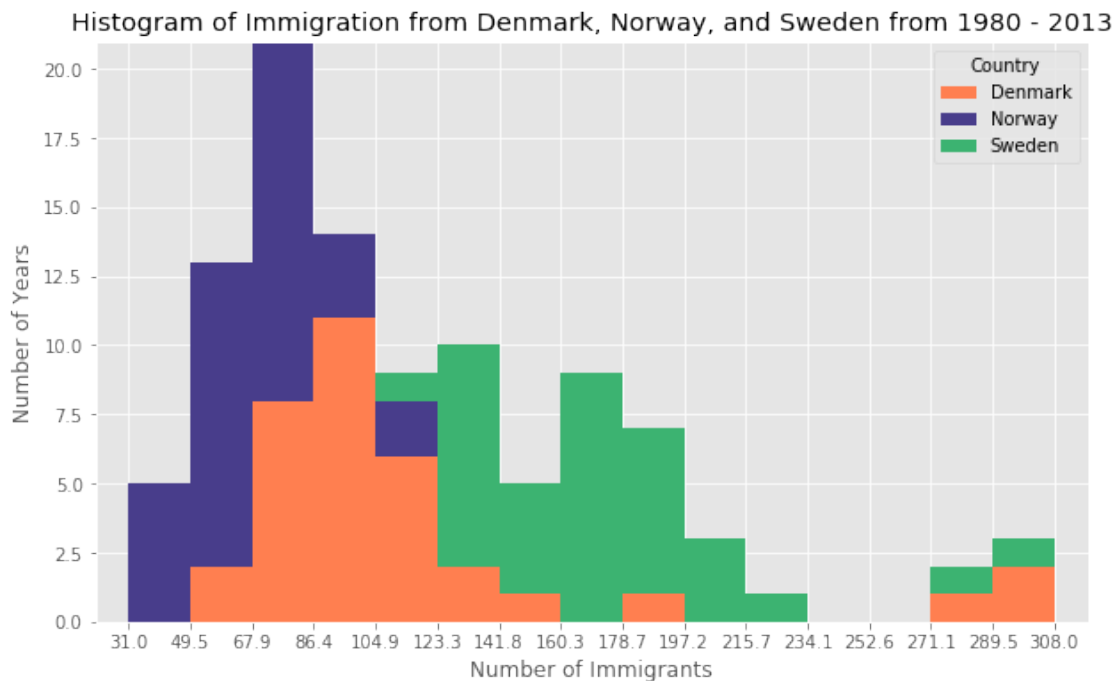
of the plot. We can pass a tuple (min,max) using the `xlim` parameter, as show below.

```
In [31]: count, bin_edges = np.histogram(df_t, 15)
        xmin = bin_edges[0] - 10 # first bin value is 31.0, adding buffer of 10 for aestheti
        xmax = bin_edges[-1] + 10 # last bin value is 308.0, adding buffer of 10 for aestheti

        # stacked Histogram
        df_t.plot(kind='hist',
                  figsize=(10, 6),
                  bins=15,
                  xticks=bin_edges,
                  color=['coral', 'darkslateblue', 'mediumseagreen'],
                  stacked=True,
                  xlim=(xmin, xmax)
        )

        plt.title('Histogram of Immigration from Denmark, Norway, and Sweden from 1980 - 2013')
        plt.ylabel('Number of Years')
        plt.xlabel('Number of Immigrants')

        plt.show()
```



Question: Use the scripting layer to display the immigration distribution for Greece, Albania, and Bulgaria for years 1980 - 2013? Use an overlapping plot with 15 bins and a transparency value of 0.35.

```

In [32]: ### type your answer here
df_cof = df_can.loc[['Greece', 'Albania', 'Bulgaria'], years]

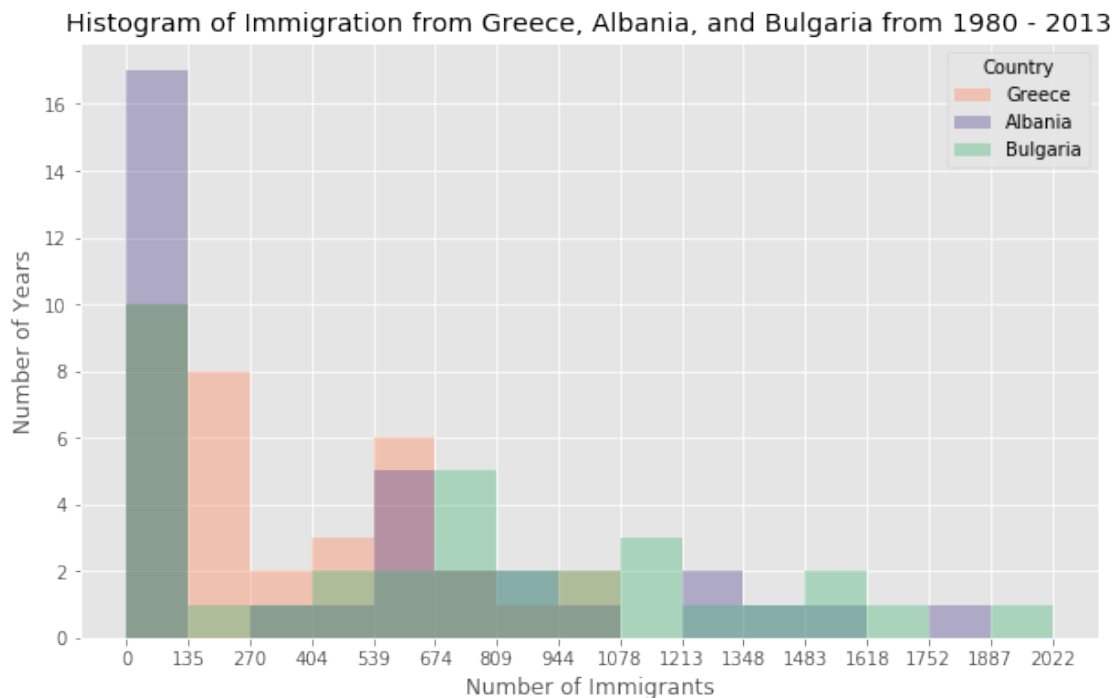
# transpose the dataframe
df_cof = df_cof.transpose()

# let's get the x-tick values
count, bin_edges = np.histogram(df_cof, 15)

# Un-stacked Histogram
df_cof.plot(kind='hist',
            figsize=(10, 6),
            bins=15,
            alpha=0.35,
            xticks=bin_edges,
            color=['coral', 'darkslateblue', 'mediumseagreen']
            )

plt.title('Histogram of Immigration from Greece, Albania, and Bulgaria from 1980 - 2013')
plt.ylabel('Number of Years')
plt.xlabel('Number of Immigrants')
plt.show()

```



6 Bar Charts (Dataframe)

A bar plot is a way of representing data where the *length* of the bars represents the magnitude/size of the feature/variable. Bar graphs usually represent numerical and categorical variables grouped in intervals.

To create a bar plot, we can pass one of two arguments via `kind` parameter in `plot()`:

- `kind=bar` creates a *vertical* bar plot
- `kind=barh` creates a *horizontal* bar plot

Vertical bar plot

In vertical bar graphs, the x-axis is used for labelling, and the length of bars on the y-axis corresponds to the magnitude of the variable being measured. Vertical bar graphs are particularly useful in analyzing time series data. One disadvantage is that they lack space for text labelling at the foot of each bar.

Let's start off by analyzing the effect of Iceland's Financial Crisis:

The 2008 - 2011 Icelandic Financial Crisis was a major economic and political event in Iceland. Relative to the size of its economy, Iceland's systemic banking collapse was the largest experienced by any country in economic history. The crisis led to a severe economic depression in 2008 - 2011 and significant political unrest.

Question: Let's compare the number of Icelandic immigrants (country = 'Iceland') to Canada from year 1980 to 2013.

```
In [33]: # step 1: get the data
```

```
df_iceland = df_can.loc['Iceland', years]
df_iceland.head()
```

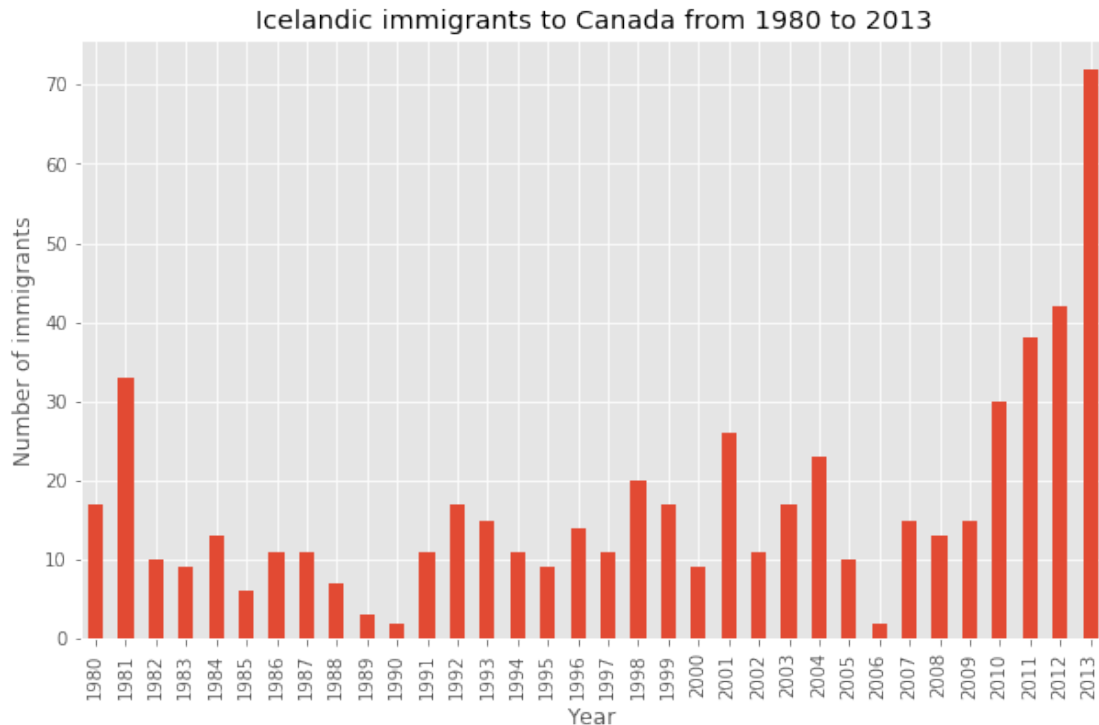
```
Out[33]: 1980    17
         1981    33
         1982    10
         1983     9
         1984    13
         Name: Iceland, dtype: object
```

```
In [34]: # step 2: plot data
```

```
df_iceland.plot(kind='bar', figsize=(10, 6))

plt.xlabel('Year') # add to x-label to the plot
plt.ylabel('Number of immigrants') # add y-label to the plot
plt.title('Icelandic immigrants to Canada from 1980 to 2013') # add title to the plot

plt.show()
```



The bar plot above shows the total number of immigrants broken down by each year. We can clearly see the impact of the financial crisis; the number of immigrants to Canada started increasing rapidly after 2008.

Let's annotate this on the plot using the `annotate` method of the **scripting layer** or the **pyplot interface**. We will pass in the following parameters: - `s`: str, the text of annotation. - `xy`: Tuple specifying the (x,y) point to annotate (in this case, end point of arrow). - `xytext`: Tuple specifying the (x,y) point to place the text (in this case, start point of arrow). - `xycoords`: The coordinate system that `xy` is given in - 'data' uses the coordinate system of the object being annotated (default). - `arrowprops`: Takes a dictionary of properties to draw the arrow: - `arrowstyle`: Specifies the arrow style, '`->`' is standard arrow. - `connectionstyle`: Specifies the connection type. `arc3` is a straight line. - `color`: Specifies color of arrow. - `lw`: Specifies the line width.

I encourage you to read the Matplotlib documentation for more details on annotations: http://matplotlib.org/api/pyplot_api.html#matplotlib.pyplot.annotate.

```
In [35]: df_iceland.plot(kind='bar', figsize=(10, 6), rot=90) # rotate the bars by 90 degrees

plt.xlabel('Year')
plt.ylabel('Number of Immigrants')
plt.title('Icelandic Immigrants to Canada from 1980 to 2013')

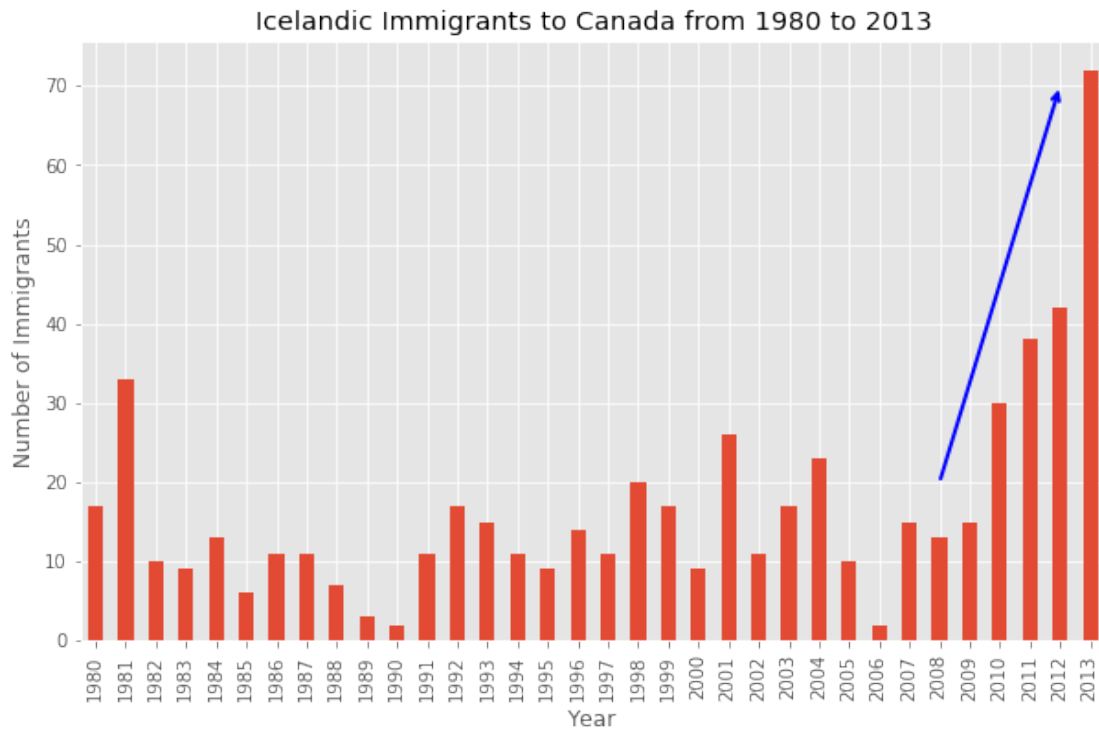
# Annotate arrow
plt.annotate('',
             xy=(32, 70),
             xytext=(28, 20),
             # s: str. Will leave it blank for no text
             # place head of the arrow at point (year 2012 , p
             # place base of the arrow at point (year 2008 , p
```

```

xycoords='data',          # will use the coordinate system of the object being annotated
arrowprops=dict(arrowstyle='->', connectionstyle='arc3', color='blue', lw=2)
)

```

```
plt.show()
```



Let's also annotate a text to go over the arrow. We will pass in the following additional parameters: - rotation: rotation angle of text in degrees (counter clockwise) - va: vertical alignment of text ['center' | 'top' | 'bottom' | 'baseline'] - ha: horizontal alignment of text ['center' | 'right' | 'left']

```
In [36]: df_iceland.plot(kind='bar', figsize=(10, 6), rot=90)
```

```

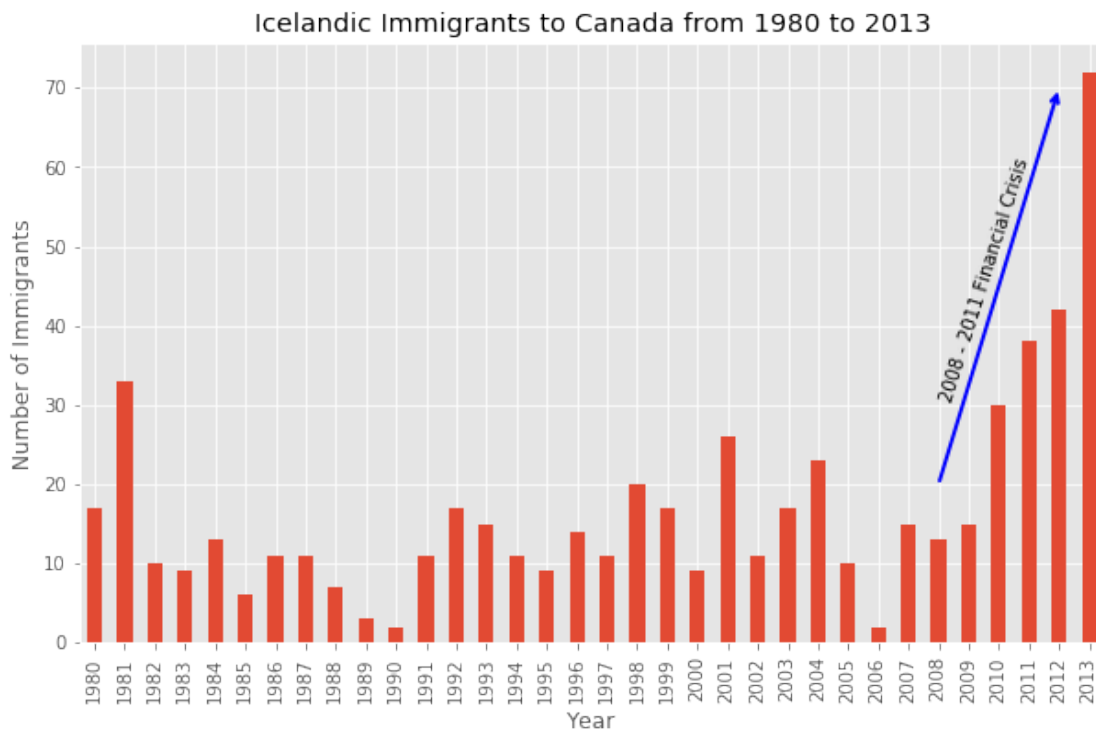
plt.xlabel('Year')
plt.ylabel('Number of Immigrants')
plt.title('Icelandic Immigrants to Canada from 1980 to 2013')

# Annotate arrow
plt.annotate('',          # s: str. will leave it blank for no text
            xy=(32, 70),  # place head of the arrow at point (year 2012 , p
            xytext=(28, 20), # place base of the arrow at point (year 2008 , p
            xycoords='data', # will use the coordinate system of the object being annotated
            arrowprops=dict(arrowstyle='->', connectionstyle='arc3', color='blue', lw=2)
)

```

```
# Annotate Text
plt.annotate('2008 - 2011 Financial Crisis', # text to display
            xy=(28, 30), # start the text at at point (year 2008 ,
            rotation=72.5, # based on trial and error to match the an
            va='bottom', # want the text to be vertically 'bottom'
            ha='left', # want the text to be horizontally 'left'
            )

plt.show()
```



Horizontal Bar Plot

Sometimes it is more practical to represent the data horizontally, especially if you need more room for labelling the bars. In horizontal bar graphs, the y-axis is used for labelling, and the length of bars on the x-axis corresponds to the magnitude of the variable being measured. As you will see, there is more room on the y-axis to label categorical variables.

Question: Using the scripting layer and the `df_can` dataset, create a *horizontal* bar plot showing the *total* number of immigrants to Canada from the top 15 countries, for the period 1980 - 2013. Label each country with the total immigrant count.

Step 1: Get the data pertaining to the top 15 countries.

```
In [37]: ### type your answer here
# sort dataframe on 'Total' column (descending)
df_can.sort_values(by='Total', ascending=True, inplace=True)
```



```
# get top 15 countries
df_top15 = df_can['Total'].tail(15)
df_top15
```

```
Out[37]: Country
Romania          93585
Viet Nam         97146
Jamaica          106431
France           109091
Lebanon          115359
Poland           139241
Republic of Korea 142581
Sri Lanka        148358
Iran (Islamic Republic of) 175923
United States of America 241122
Pakistan         241600
Philippines      511391
United Kingdom of Great Britain and Northern Ireland 551500
China            659962
India            691904
Name: Total, dtype: int64
```

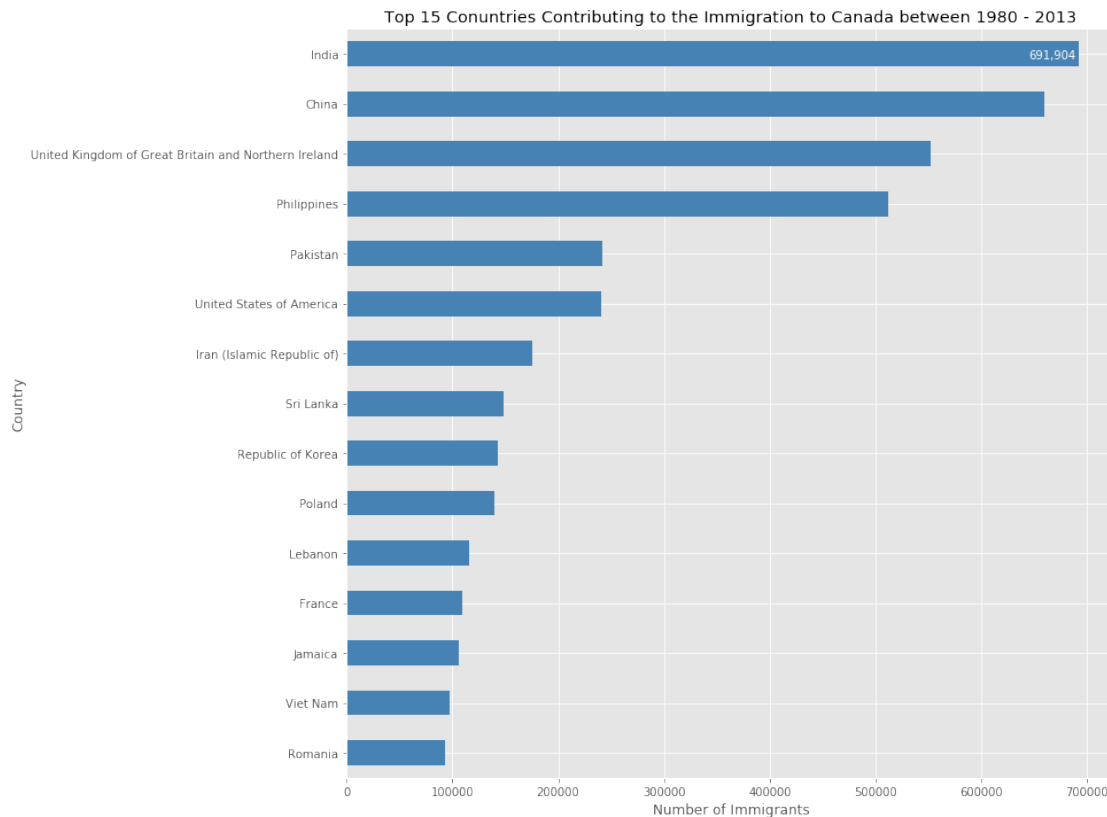
Step 2: Plot data: 1. Use kind='barh' to generate a bar chart with horizontal bars. 2. Make sure to choose a good size for the plot and to label your axes and to give the plot a title. 3. Loop through the countries and annotate the immigrant population using the anotate function of the scripting interface.

```
In [38]: ### type your answer here
df_top15.plot(kind='barh', figsize=(12, 12), color='steelblue')
plt.xlabel('Number of Immigrants')
plt.title('Top 15 Conuntries Contributing to the Immigration to Canada between 1980 - 2015')

# annotate value labels to each country
for index, value in enumerate(df_top15):
    label = format(int(value), ',') # format int with commas

# place text at the end of bar (subtracting 47000 from x, and 0.1 from y to make it fit)
plt.annotate(label, xy=(value - 47000, index - 0.10), color='white')

plt.show()
```



6.0.2 Thank you for completing this lab!

This notebook was originally created by [Jay Rajasekharan](#) with contributions from [Ehsan M. Kermani](#), and [Slobodan Markovic](#).

This notebook was recently revamped by [Alex Aklson](#). I hope you found this lab session interesting. Feel free to contact me if you have any questions!

This notebook is part of a course on **Coursera** called *Data Visualization with Python*. If you accessed this notebook outside the course, you can take this course online by clicking [here](#).

Copyright © 2018 [Cognitive Class](#). This notebook and its source code are released under the terms of the [MIT License](#).